

Fibonacci Heaps [Introduction to Algorithms by CLRS]

dual purpose:

- 1) supports a set of operations that constitutes a "mergeable heap".
- 2) several operations run in constant amortized time.

mergeable heap:

- MAKE-HEAP()
- INSERT(H, x) inserts element x , whose key has already been filled in, into heap H .
- MINIMUM(H) returns a pointer to the element in heap H whose key is minimum.
- EXTRACT-MIN(H)
- UNION(H_1, H_2)

Fibonacci heaps also support:

- + DECREASE-KEY(H, x, k) assigns to element x within H the new key value k , which we assume to be no greater than its current key value.
- + DELETE(H, x)

Procedure	Binary heap (worst-case)	Fibonacci heap (amortized)
MAKE-HEAP	$\Theta(1)$	$\Theta(1)$
INSERT	$\Theta(\lg n)$	$\Theta(1)$
MINIMUM	$\Theta(1)$	$\Theta(1)$
EXTRACT-MIN	$\Theta(\lg n)$	$\Theta(\lg n)$
UNION	$\Theta(n)$	$\Theta(1)$
DECREASE-KEY	$\Theta(\lg n)$	$\Theta(1)$
DELETE	$\Theta(\lg n)$	$\Theta(\lg n)$

desirable when number of these operations are small

Fibonacci heap:

A Fibonacci heap is a collection of rooted trees that are min-heap ordered.

min-heap property:

the key of a node is greater than or equal to the key of its parent.

application:

- 1) counting minimum spanning trees
- 2) single-source shortest paths.

drawbacks:

- large constant factors
- programming complexity

Potential function

$t(H)$ the number of trees in the root list

$m(H)$ the number of marked nodes in H .

$$\Phi(H) = t(H) + 2m(H)$$

- assume that a unit of potential can pay for a constant amount of work.

Maximum degree

assume that we know an upper bound $D(n)$ on the maximum degree of any node in an n -node Fibonacci heap.

$$D(n) \leq \lfloor \lg n \rfloor$$

when we support DECREASE-KEY and DELETE, $D(n) = O(\lg n)$.

Inserting a node

- just add it to the root list

- the increase in potential $t(H') = t(H) + 1$, $m(H') = m(H)$ is 1, actual cost $O(1)$, amortized cost $O(1) + 1 = O(1)$

Uniting two Fibonacci heaps

- change in potential

$$\begin{aligned}\Phi(H) &= (\Phi(H_1) + \Phi(H_2)) \\ &= (t(H_1) + 2m(H_1)) + (t(H_2) + 2m(H_2)) \\ &= 0\end{aligned}$$

The amortized cost of FIB-HEAP-UNION is equal to its $O(1)$ actual cost.

Extracting the minimum node

FIB-HEAP-EXTRACT-MIN(H)

$z = H.min$

if $z \neq NIL$

for each child x of z

add x to the root list of H

$x.p = NIL$

remove z from the root list of H

if $z == z.right$

$H.min = NIL$

else $H.min = z.right$ ← not necessarily going to be the new minimum node.

CONSOLIDATE(H)

$H.n = H.n - 1$

return z

CONSOLIDATE (H)

need to know upper bound.

let $A[0..D(H,n)]$ be a new array // keep track of roots according to their degrees

for $i = 0$ to $D(H,n)$

$A[i] = NIL$

for each node w in the root list of H .

$x = w$

$d = x.degree$

while $A[d] \neq NIL$

$y = A[d]$ // another node with the same degree as x

if $x.key > y.key$

exchange x with y

FIB-HEAP-LINK(H, y, x)

$A[d] = NIL$

$d = d + 1$

$A[d] = x$

$H.min = NIL$

for $i = 0$ to $D(H,n)$

if $A[i] \neq NIL$

if $H.min == NIL$

create a root list of H containing just $A[i]$

$H.min = A[i]$

else insert $A[i]$ into H 's root list

if $A[i].key < H.min.key$

$H.min = A[i]$

FIB-HEAP-LINK(H, y, x)

1. remove y from the root list of H

2. make y a child of x , incrementing $x.degree$

3. $y.mark = FALSE$

amortized cost : try to show it is $O(D(n))$.

① $O(D(n))$ contribution comes from FIB-HEAP-EXTRACT-MIN processing at most $D(n)$ children of the minimum node.

Thus, the total actual work is ①② in $O(D(n) + t(H))$.

potential before extracting minimum $t(H) + 2m(H)$

potential afterwards is at most $(D(n) + 1) + 2m(H)$

$$O(D(n) + t(H)) + ((D(n) + 1) + 2m(H)) - (t(H) + 2m(H)) = O(D(n)) + O(t(H)) - t(H) = O(D(n))$$

can scale up the units of potential to dominate the constant.

loop invariant:

At the start of each iteration of the while loop, $d = x.degree$

② the total amount of work in the for loop is at most proportional to $D(n) + t(H)$

Bounding the maximum degree

to show the upper bound of $D(n)$ is $O(\lg n)$.

In particular, $D(n) \leq \lfloor \log_{\phi} n \rfloor$

$$\phi = \frac{1+\sqrt{5}}{2}$$

Lemma 19.1

Let x be any node in a Fibonacci heap, and suppose that $x.\text{degree} = k$. Let y_1, y_2, \dots, y_k denote the children of x in the order in which they were linked to x from the earliest to the latest. Then $y_1.\text{degree} \geq 0$ and $y_i.\text{degree} \geq i-2$ for $i=2, 3, \dots, k$.

Lemma 19.2

For all integers $k \geq 0$,

$$F_{k+2} = 1 + \sum_{i=0}^k F_i$$

$$F_k = \begin{cases} 0 & k=0 \\ 1 & k=1 \\ F_{k-1} + F_{k-2} & k \geq 2 \end{cases}$$

Lemma 19.3

For all integers $k \geq 0$, the $(k+2)$ nd Fibonacci number satisfies $F_{k+2} \geq \phi^k$

inductive step:

$$\begin{aligned} F_{k+2} &= F_{k+1} + F_k \\ &\geq \phi^{k-1} + \phi^{k-2} \quad (\text{by the inductive hypothesis}) \\ &= \phi^{k-2} (\phi + 1) \\ &= \phi^{k-2} \cdot \phi^2 \quad (\phi \text{ is the positive root of equation } x^2 = x + 1) \\ &= \phi^k \end{aligned}$$

Lemma 19.4

Let x be any node in a Fibonacci heap, and let $k = x.\text{degree}$. Then $\text{size}(x) \geq F_{k+1}$

$$F_{k+1} \geq \phi^k, \text{ where } \phi = \frac{1+\sqrt{5}}{2}$$

Proof. Let S_k denote minimum possible size of any node of degree k in any Fibonacci heap.

$$S_k \leq \text{size}(x)$$

$$\text{size}(x) \geq S_k$$

$$\geq 2 + \sum_{i=2}^k S_{y_i.\text{degree}}$$

$$\geq 2 + \sum_{i=2}^k S_{i-2}$$

$$\geq 2 + \sum_{i=0}^k F_i$$

$$= 1 + \sum_{i=0}^k F_i$$

$$= F_{k+2} \quad (\text{by Lemma 19.2})$$

$$\geq \phi^k \quad (\text{by Lemma 19.3})$$

$$\therefore \text{size}(x) \geq S_k \geq F_{k+2} \geq \phi^k$$

Corollary 19.5

The maximum degree $D(n)$ of any node in an n -node Fibonacci heap is $O(\lg n)$.

$$n \geq \text{size}(x) \geq \phi^k \text{ where } k = x.\text{degree}$$

$$k \leq \lfloor \log_{\phi} n \rfloor$$

Maximum Flows and Parametric Shortest Paths in Planar Graphs

by Jeff Erickson jeffe@cs.uiuc.edu

let $G = (V, E)$ be a directed plane graph. s, t be vertices of G
let $c: E \rightarrow \mathbb{R}$ be a nonnegative capacity function

Goal: compute a maximum (s, t) -flow in G .

Assume WLOG the reversal of any directed edge in G is also an edge in G
 \longrightarrow implies that both G and its dual G^* are strongly connected.

Venkatesan's Reduction

Idea: compute a feasible (s, t) -flow with fixed value λ , or correctly report that no such flow exists, by reduction to a single source shortest path problem in an appropriately weighted dual graph G^* .

$\pi(e)$: Fixed an arbitrary directed path P from s to t , and let $\pi: E \rightarrow \mathbb{R}$ denote the unit flow through P :

$$\pi(e) := \begin{cases} 1 & \text{if } e \in P \\ -1 & \text{if } \text{rev}(e) \in P \\ 0 & \text{otherwise} \end{cases}$$


For any subset $E' \subseteq E$, let $\pi(E') = \sum_{e \in E'} \pi(e)$.

cocycle: A subgraph C of G is called a cocycle if the corresponding dual subgraph C^* is a simple directed cycle in G^* .

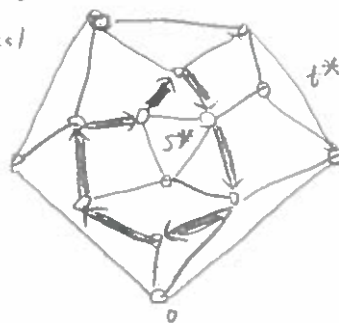
crossing number: For any cycle C^* in G^* , we call $\pi(C)$ the crossing number of C^* .

Lemma 2.1: $\pi(C) \in \{-1, 0, 1\}$ for any cocycle C . Moreover, $\pi(C) = 1$ if and only if C is an (s, t) -cut.

original:



dual



consider the flow $\lambda \cdot \pi$, which assigns value λ to every directed edge in path P
 $\begin{cases} -\lambda & \text{to every edge in } \text{rev}(P) \\ 0 & \text{to every other edge.} \end{cases}$

Let $G_\lambda := G_{\lambda, \pi}$ denote the residual network of this flow

capacity function: $c(\lambda, e) := c(e) - \lambda \cdot \pi(e)$.

$\therefore \lambda \cdot \pi$ is feasible if and only if $c(\lambda, e) \geq 0$ for every edge e in G .

Let G_λ^* denote the dual residual network, which is the directed dual graph G^* where every edge e^* has a cost $c(\lambda, e^*) = c(\lambda, e)$.

Lemma 2.2 There is a feasible (s, t) -flow in G with value λ if and only if the residual network G_λ^* does not contain a negative cycle.

let $\text{dist}(\lambda, p)$ denote the shortest path distance in G_λ^* from o to p . (arbitrary dual vertex, called origin)

define $\phi(\lambda, e) := \text{dist}(\lambda, \text{head}(e^*)) - \text{dist}(\lambda, \text{tail}(e^*)) + \lambda \cdot \pi(e)$.

because the duals of the edges leaving v define a directed cycle in G^* all the $\text{dist}(\lambda, \cdot)$ terms in the sum cancel out.

$\therefore \phi(\lambda, \cdot)$ is a valid (s, t) -flow with value λ .

define the slack of each dual edge e^* :

$$\text{slack}(\lambda, e^*) := \text{dist}(\lambda, \text{tail}(e^*)) + c(\lambda, e) - \text{dist}(\lambda, \text{head}(e^*))$$

$$\text{so } \text{slack}(\lambda, e^*) = c(e) - \phi(\lambda, e)$$

Parametric Shortest Paths

Let λ_{\max} denote the largest value of λ for which shortest paths in G_λ^* are well-defined. Lemma 2.2 implies that λ_{\max} is also the value of the maximum flow.

For any particular value of λ , let T_λ denote the single-source shortest path tree in G_λ^* rooted at o .

High level algorithm

PLANAR MAX FLOW (G, c, s, t):

Compute T .

Maintain T_λ as λ increases continuously from 0

to λ_{\max} .
 Compute $\phi(\lambda_{\max}, \cdot)$ from $T_{\lambda_{\max}}$

Genericity assumption

we assume that the capacity function is generic.

- 1) Our genericity assumption implies that T_λ is uniquely defined for all λ between 0 and λ_{\max} , except for a finite set of critical values
- 2) Our genericity assumption implies that exactly one non-tree edge becomes tense at each critical value of λ .

Lemma 2.3 λ_{\max} is the first critical value of λ whose pivot introduces a directed cycle into T_λ .

pivot At each critical value of λ , some non-tree dual edge $p \rightarrow q$ becomes tense and enters T_λ , replacing the previous edge $\text{pred}(\lambda, q) \rightarrow q$.

tense call a dual edge e^* tense if $\text{slack}(\lambda, e^*) = 0$.

Lemma 2.4 λ_{\max} is the smallest critical value of λ whose pivot disconnects L_λ

loose call a primal edge e loose at λ if neither its dual e^* nor its reversed dual $\text{rev}(e^*)$ is tense at λ , and let L_λ be the graph of all loose edges.

active A dual edge is active at λ if its slack at λ is decreasing

LP_λ The primal spanning tree L_λ contains a unique directed path from s to t ; call this loose path LP_λ .

Lemma 2.5 A dual edge e^* is active at λ if and only if e is an edge of LP_λ .

★ PLANAR MAX FLOW (G, c, s, t) :

Initialize the spanning tree L , predecessors and slacks
while s and t are in the same component of L

$LP \leftarrow$ the path from s to t

$p \rightarrow q \leftarrow$ the edge in LP^* with minimum slack

$\Delta \leftarrow \text{slack}(p \rightarrow q)$

for every edge e in LP

$\text{slack}(e^*) \leftarrow \text{slack}(e^*) - \Delta$

$\text{slack}(\text{rev}(e^*)) \leftarrow \text{slack}(\text{rev}(e^*)) + \Delta$

delete $(p \rightarrow q)^*$ from L

if $q \neq 0$ << that is, if $\text{pred}(q) \neq \emptyset$ >>

insert $(\text{pred}(q) \rightarrow q)^*$ into L

$\text{pred}(q) \leftarrow p$

for each edge e

$\phi(e) \leftarrow c(e) - \text{slack}(e^*)$

} mutates T_λ as λ
increases by pivoting
until λ_{\max}

} compute the flow